

# practical\_ml

August 6, 2019

## 0.1 Practical Machine Learning Assignment

```
In [1]: from sklearn.multiclass import OneVsOneClassifier, OneVsRestClassifier
        from sklearn.metrics import classification_report
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.model_selection import cross_val_score
        from sklearn.preprocessing import LabelEncoder
        from collections import defaultdict
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
```

## 0.2 Summary of Prediction

### Note: I used Python, not R

In this project, I used a Decision Tree classifier into a one-vs-one multiclass classifier. The reason I chose a Tree-based model is because of the large number of feature in the data. Using a linear model such as Logistic Regression or Support Vector Machine would have required one-hot encoding which would have increased the dimension even further. We do not have that problem with Tree-Based models, as label encoding works just fine.

I started simple with a Decision Tree and it turned out to be the best model. I found that random forest overfitted the train set even with a small number of Trees in the forest. Also, I had to drop many features which had more than 50% missing values, which reduced the dimension from 160 to 60.

I used 5-fold cross validation to estimate the accuracy on the test set. This gave me a value of 0.53.

### 0.2.1 Data Import and Exploration

```
In [2]: df = pd.read_csv('pml-training.csv', low_memory=False)
```

```
In [3]: print(df.shape)
        df[df.columns[0:10]].head()
```

```
(19622, 160)
```

```
Out [3]: Unnamed: 0 user_name raw_timestamp_part_1 raw_timestamp_part_2 \
0          1 carlitos          1323084231          788290
1          2 carlitos          1323084231          808298
2          3 carlitos          1323084231          820366
3          4 carlitos          1323084232          120339
4          5 carlitos          1323084232          196328
```

	cvtd_timestamp	new_window	num_window	roll_belt	pitch_belt	yaw_belt
0	05/12/2011 11:23	no	11	1.41	8.07	-94.4
1	05/12/2011 11:23	no	11	1.41	8.07	-94.4
2	05/12/2011 11:23	no	11	1.42	8.07	-94.4
3	05/12/2011 11:23	no	12	1.48	8.05	-94.4
4	05/12/2011 11:23	no	12	1.48	8.07	-94.4

```
In [4]: y = df['classe']
df_clean = df.drop(['Unnamed: 0', 'classe'], axis=1)
```

```
In [5]: #find columns with large number of missing values
missing = []
for column in df_clean.columns:
    if df_clean[column].isna().sum()/df_clean.shape[0] > 0.5:
        missing.append(column)
df_clean = df_clean.drop(missing, axis=1)

objects = df_clean.select_dtypes(include=['object'])
numerics = df_clean.select_dtypes(include=['int', 'float64', 'int64'])
print(len(objects.columns) + len(numerics.columns) == len(df_clean.columns))
print(objects.shape, numerics.shape)
```

```
True
(19622, 3) (19622, 55)
```

```
In [6]: d = defaultdict(LabelEncoder)
categoric = objects.apply(lambda x: d[x.name].fit_transform(x))
```

```
In [7]: X_train = pd.concat([categoric, numerics], axis=1, sort=False)
y_train = LabelEncoder().fit_transform(y)
```

```
In [8]: predictor = OneVsOneClassifier(DecisionTreeClassifier(random_state = 1, max_depth = 5))
y_pred = predictor.fit(X_train, y_train).predict(X_train)
target_names = ['class 0', 'class 1', 'class 2', 'class 3', 'class 4']
print(classification_report(y_train, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
class 0	0.93	0.89	0.91	5580
class 1	0.90	0.78	0.83	3797
class 2	0.86	0.92	0.89	3422

class 3	0.78	0.90	0.83	3216
class 4	0.93	0.94	0.93	3607
micro avg	0.88	0.88	0.88	19622
macro avg	0.88	0.88	0.88	19622
weighted avg	0.89	0.88	0.88	19622

## 0.2.2 Cross-Validation Score

```
In [9]: scores = cross_val_score(predictor, X_train, y_train, cv = 5)
        print("Mean Cross-Validation Score: ", np.mean(scores))
```

Mean Cross-Validation Score: 0.5266768706920635

## 0.2.3 Test Set Performance

```
In [10]: df_test = pd.read_csv('pml-testing.csv', low_memory=False).drop(['Unnamed: 0'], axis=0)
        print(df_test.shape)
        df_test = df_test.drop(missing, axis=1)
        print(df_test.shape)
        objects_df = df_test.select_dtypes(include=['object'])
        numerics_df = df_test.select_dtypes(include=['int', 'float64', 'int64'])
        print(objects_df.shape, numerics_df.shape)
        print(len(objects_df.columns) + len(numerics_df.columns) == len(df_test.columns))

        categoric_df = objects_df.apply(lambda x: d[x.name].transform(x))
        X_test = pd.concat([categoric_df, numerics_df], axis=1, sort=False)[X_train.columns]

        y_test = predictor.predict(X_test)
```

```
(20, 159)
(20, 59)
(20, 3) (20, 56)
True
```

```
In [11]: y_test
```

```
Out[11]: array([1, 0, 2, 0, 0, 4, 3, 3, 0, 0, 1, 2, 1, 0, 4, 4, 0, 3, 1, 1])
```

```
In [12]: y.unique()
```

```
Out[12]: array(['A', 'B', 'C', 'D', 'E'], dtype=object)
```